

Memnon

A Local-First Voice Pipeline

Voice is the fastest way to capture an idea. But recordings rot.


- You record something valuable — it sits in Voice Memos, never found again
- Plaud AI, Pocket AI, Limitless: hardware (\$160+), subscriptions, your audio to their cloud
- Apple Voice Memos transcription (iOS 18): plain text, no structure, no knowledge base
- whisper.cpp exists — but an iPhone-first, Obsidian-native pipeline you can fork is still the gap



"Do we really need proprietary hardware and subscriptions just to use our own voice notes?"

Local transcription already exists.

whisper.cpp runs on your Mac for free.



- Audio transcribed locally — never leaves your machine
- AI summarization optional — works without it
- Output is plain Markdown — universal, readable, future-proof
- One Python file, minimal dependencies (edge-tts and ffmpeg only for audio output)

This is where it gets honest. And interesting.

The dream: record on iPhone, note appears in Obsidian. **The reality: Apple controls the capture layer.**

Attempt 1 — Custom Shortcuts Works, but screen must stay on. Lock your phone and recording stops.

Attempt 2 — Voice Memos automation "Get Latest Voice Memo" action. Works on some devices. Simply absent on others — Apple won't say why.

Attempt 3 — Manual handoff (current) Record → Stop → Share → Save to Files → Voice Inbox/raw. One step. Reliable everywhere.

Layer	Who controls it	Stability
Mac pipeline	You	Completely stable
iPhone capture	Apple	iOS version / device dependent

"Stop trying to win on the iPhone side. Make the Mac side excellent."

The Action Button changes the equation.

- iPhone 15 Pro + all iPhone 16 models
- Configure to open Voice Memos directly
- Press → record → lock screen → stop when done
- One manual step: Share → Save to Files → raw

This is what makes Memnon a real Plaud / Pocket AI alternative:

- ✓ No dedicated hardware
- ✓ No subscription
- ✓ No audio to the cloud

Action Button



*Uses Apple's strongest native
recording surface*

TypeScript wasn't in the original design. It earned its place.

After the Python pipeline was stable, the question became: how do you query your knowledge base?

Built a TypeScript MCP server that exposes Obsidian notes to Claude

list_notes

search_notes

get_note

get_action_items

"What action items do I have from this week?"

"Search my notes for anything about the dog park idea"

"Summarize what I've been thinking about this month"

v1.0.0-python — pure Python

v2.0.0-mcp — adds TypeScript MCP

Started as transcription. Became a personal AI knowledge pipeline.



Capture (iPhone)



Raw ingest (iCloud Drive)



Transcription (whisper.cpp, local)



Lane detection (voice label → heuristics → default)



Lane-aware AI summarization



Obsidian note (Markdown)



Lane actions (per-lane downstream pipelines)



MCP server (Claude can query everything)



Different recordings deserve different treatment.

Professional: just met Sarah who works in EdTech...

Build: thinking about the API redesign...

Reflect: feeling stuck on the same pattern...

Good Morning: what should I focus on today?

Falls back to keyword heuristics, then default.

Lane	Output
build	Voice note: title, summary, action items, tags
professional	Voice note + AI coaching note: observations, next steps, to-do checklist
reflect	Voice note + wisdom note + podcast MP3 + meditation MP3
good morning	Energizing briefing + reflection synthesis + combined daily podcast

The pipeline produces Markdown.

But it also produces audio.

reflect lane

- Records a 2-minute voice note
- AI selects 3–5 passages from a private corpus
- 215 passages · 29 authors: Plato, Kierkegaard, Psalms, Pascal, Heidegger...
- Generates podcast + meditation scripts
- edge-tts synthesis (free, no API key)
- Mood-based ambient music via ffmpeg
- 3 files land in iCloud: note · podcast · meditation

good morning lane

- Say "Good Morning:" into your phone
- Reads last 7 days of wisdom notes + professional context
- Energizing briefing: what matters today
- Weekly reflection synthesis
- Musical crossfade between segments
- Waiting in iCloud by the time you've had coffee

"This is what 'your downstream workflows' actually looks like once you own the pipeline."

The repo is the pipeline.

Your config is the workflow.

Core repo owns	Your private config owns
Raw ingest	Personal prompts
Transcription	GPT connections
Lane routing framework	Lane definitions
Generic templates	Corpus of texts
MCP server	Audio output directories
Status artifact (last-run.json)	Calendar/CRM wiring

"The pipeline outputs plain Markdown to Obsidian. Every companion tool consumes from there. No tool should need to change the pipeline to add a feature — if it does, the boundary is in the wrong place."

Live Demo

Option A: Voice note (live)

1. Press Action Button → Voice Memos opens
2. Say "Professional:" then speak naturally
3. Stop → Share → Save to Files → Voice Inbox → raw
4. Wait ~60 seconds
5. Note in Obsidian: title, summary, coaching note, to-do
6. Ask Claude: "What did I just say?"

Option B: Daily podcast (pre-recorded)

Play demo-daily-bouldersjs.mp3



"Generated this morning from a real voice note."

A reference architecture for local-first AI workflows.

- Small enough to read in an afternoon
- Designed to be forked and extended
- No always-on server, no vendor lock-in
- Composable — swap the transcriber, AI backend, note format, destination
- The audio layer is entirely optional (edge-tts, ffmpeg)

"The pipeline is stupid about where audio comes from. That's a feature."

A competitor gives you their format or nothing. Memnon gives you their format and everything else.